

The Future Of Machine Learning And Artificial Intelligence: Key Trends And Innovative Applications

Sunil Kumar Sahu¹, Abhay Kumar², Arjit Singh³, Pramveer Kumar⁴

¹Assistant Professor, Computer Science and Engineering Supaul College of Engineering, Supaul
Email id: sunil.sahu01@bihar.gov.in ORCID ID: 0009-0002-2125-0298

²Assistant Professor, Computer Science and Engineering Supaul College of Engineering, Supaul
Email id: abhaybvpit@gmail.com ORCID ID: 0000-0002-0348-1562

³B.Tech (Student), Computer Science and Engineering Supaul College of Engineering, Supaul-852131
Email id: arjit.codes04@gmail.com ORCID ID: 0009-0000-9700-1760

⁴B.Tech (Student), Computer Science and Engineering (Artificial Intelligence) Supaul College of Engineering, Supaul Email id: pramveerkumar9939@gmail.com
ORCID ID: 0009-0002-7729-9314

Abstract

This paper examines near-term trajectories in machine learning and artificial intelligence (2025–2030), synthesizing advances in multimodal foundation models, on-device small models, agentic tool use, synthetic data, and specialized hardware with evolving governance frameworks. Using a narrative review with systematic elements across peer-reviewed studies, benchmarks, and policy documents, we identify five trends that will shape deployment: trustworthy multimodality, privacy-preserving personalization, energy-aware efficiency, rigorous evaluation beyond static benchmarks, and integration of AI agents into real workflows. We map these trends to innovative applications in life sciences, education, finance, climate and energy, manufacturing, agriculture, and public services, highlighting measurable benefits (e.g., faster discovery, improved forecasting, and productivity gains) alongside risks related to safety, bias, provenance, and compliance. The paper concludes with a research agenda prioritizing open evaluation ecosystems, verifiable data governance, and hybrid cloud/edge architectures. Findings provide practitioners and policymakers with actionable guidance for scaling capable, trustworthy AI through 2030 worldwide and responsibly.

Keywords: multimodal models; on-device AI; agentic workflows; synthetic data; AI governance

1. Introduction

Artificial intelligence (AI) and machine learning (ML) have moved from narrow, task-specific systems to broadly capable foundation models that reason across text, images, audio, code, and tools. This inflection is measurable. The *AI Index 2025* reports that the inference cost of achieving GPT-3.5-level performance fell by $>280\times$ between November 2022 and October 2024, while hardware costs declined $\sim 30\%$ per year and energy efficiency improved $\sim 40\%$ annually. In parallel, open-weight models have rapidly closed the performance gap with closed models on standard evaluations (from $\sim 8\%$ to $\sim 1.7\%$ on select benchmarks in a single year), expanding access for researchers and enterprises. These dynamics cheaper inference, better hardware efficiency, and

more customizable model supply explain why AI capabilities are diffusing into mainstream products and public services at an unprecedented pace.

Adoption and economic signals reinforce the technological shift. Across the OECD, enterprise use of AI surged in 2024: the EU-27 recorded 13.5% of firms (10+ employees) using AI about 60% higher than the prior year while large-firm adoption outpaces small-firm use by a factor of roughly 3.3 (39% vs. 12%), highlighting a widening SME adoption gap. Meanwhile, governments and investors are building out AI infrastructure; greenfield FDI in data centres reached \$144 billion in 2024, about 9% of global FDI, underscoring AI's role as a generalpurpose technology that depends on compute, energy, and connectivity. Although estimates vary, leading analyses suggest that generative AI could add \$2.6–\$4.4 trillion in annual value across industries, with early impact concentrated in customer operations, marketing and sales, software engineering, and R&D. Together, these data points frame AI not merely as a research frontier but as a macro-economic and organizational phenomenon.

Hardware and systems advances help explain the rapid cost/performance curve. NVIDIA's Blackwell architecture introduces a second-generation Transformer Engine and support for ultralow-precision formats (e.g., FP4), alongside reliability and serviceability features designed for massive training and long-running inference. These innovations enable higher throughput and lower energy per token for both dense and Mixture-of-Experts models practical levers for scaling agentic, multimodal systems without linearly scaling cost and carbon. As models become nimbler, small/on-device models and hybrid (device + cloud) orchestration are emerging, promising privacy, latency, and offline robustness for everyday applications.

At the same time, policy and assurance frameworks are maturing. The EU AI Act entered into force with phased obligations: bans on “unacceptable-risk” systems apply from February 2, 2025, transparency obligations for general-purpose AI (GPAI) start August 2, 2025, and high-risk requirements follow later (36 months after entry into force). The European Commission has also issued GPAI guidance to clarify scope and expectations. In the United States, NIST released the Generative AI Profile (AI 600-1) in July 2024 as a companion to the AI Risk Management Framework, offering concrete controls for data provenance, secure model operations, evaluation, and post-deployment monitoring. These timelines and profiles signal a new default: innovation and deployment must proceed with robust governance, documentation, and evaluation.

Scientific advances further justify the strategic attention on AI. Beyond language, models are increasingly science-capable: AlphaFold 3 (Nature, 2024) expanded structural prediction from single proteins to biomolecular complexes (proteins, nucleic acids, ligands, ions), enabling richer in-silico pipelines for drug discovery and synthetic biology. Such results illustrate AI's evolving role as a discovery engine, compressing cycles from hypothesis to validation and hinting at cross-domain spillovers from materials to climate modeling.

Against this backdrop, this paper takes stock of key trends multimodality and agentic tool-use, efficient on-device intelligence, synthetic data and evaluation, specialized hardware/infra, and governance and maps them to innovative applications likely to dominate the next five years (2025–2030). Our aim in the remaining sections is to synthesize the strongest empirical evidence and policy guidance, identify open problems (evaluation, energy, privacy, and safety), and outline a practical research agenda for academia, industry, and the public sector.

2. Literature Review

Foundational advances in model architecture, scaling, and training have reshaped the AI landscape over the past decade. The Transformer (Vaswani et al., 2017) displaced recurrent and convolutional sequence models by relying entirely on self-attention, enabling unprecedented parallelism and transfer across tasks. Subsequent scaling law studies argued that performance is not a simple function of parameter count alone but of a compute-optimal balance between parameters and training tokens: Hoffmann et al. (2022) showed that *Chinchilla* a 70B-parameter model trained on $\sim 4\times$ more tokens than prior practice outperformed larger but under-trained systems across a range of benchmarks, motivating today's emphasis on data quantity/quality and longer training regimes. In parallel, Mixture-of-Experts (MoE) methods (e.g., Switch Transformer) revived conditional computation: sparse expert routing allows models with “trillion-scale” parameters to retain near-constant FLOPs per token, improving throughput and latency without sacrificing accuracy (Fedus et al., 2021; see also recent MoE surveys consolidating design patterns and stability strategies). Together, these lines of work transformerization, compute-optimal scaling, and sparse routing form the bedrock of current “foundation model” capabilities and explain why frontier systems now generalize across modalities and tasks while remaining economically deployable at scale.

A second arc in the literature investigates efficiency how to compress or serve large models without measurable losses. Quantization-centric work such as LLM.int8() demonstrates that vector-wise INT8 matrix multiplication with mixed-precision “outlier” handling can cut inference memory $\sim 50\%$ while preserving quality, enabling single-node inference for 100B+ models; GPTQ further pushes *post-training* weight quantization to 3–4 bits with negligible degradation and material speedups; and SmoothQuant extends accuracy-preserving INT8 to *both* weights and activations (W8A8), smoothing activation outliers to avoid the usual accuracy cliff (Dettmers et al., 2022; Frantar et al., 2022; Xiao et al., 2023). Complementing compression, the small/edge-model literature (TinyML; on-device LMs) argues that capable models can reside on phones and microcontrollers for privacy-preserving, low-latency applications; Microsoft's Phi-3 technical report is emblematic, showing a 3.8B-parameter model trained on 3.3T tokens achieving $\sim 69\%$ MMLU while remaining deployable on consumer hardware (Abdin et al., 2024). Survey work in TinyML (ACM Computing Surveys) and monographs in federated learning (Foundations & Trends in ML) provide the systems and privacy context for these edge deployments, including communication-efficient training, personalization, and robustness under non-IID data.

Concurrently, capability expansions have come from multimodality, tool use, and agentic workflows. Visual-language models (VLMs) such as Flamingo integrated interleaved image/video

and text to deliver strong few-shot performance; the Segment Anything project released SA-1B (1 billion masks across 11 million images) and a promptable segmentation model exhibiting broad zero-shot transfer a landmark for vision foundation models and datacentric scaling (Alayrac et al., 2022; Kirillov et al., 2023). On the reasoning-and-action axis, ReAct interleaves chain-of-thought with external actions (e.g., search) to reduce hallucinations and improve decision-making, while Toolformer shows self-supervised acquisition of API-calling behavior, allowing language models to *teach themselves* to use calculators, search, and translation tools (Yao et al., 2022; Schick et al., 2023). These developments coincide with a maturing evaluation literature: MMLU standardized breadth testing across 57 academic/professional tasks; HELM advocates multi-metric, scenario-diverse evaluation (accuracy, calibration, robustness, toxicity, efficiency) with transparent prompts/completions and a “living benchmark” approach. Recent AI Index chapters document rapid gains on newly introduced hard benchmarks and call for evaluations that reflect tool use and multimodal agents rather than static Q&A alone.

Finally, the intersection of AI and science/medicine and governance has become a focal point in top venues. In biology, AlphaFold 3 (Nature, 2024) generalizes structure prediction to *biomolecular complexes* (proteins, nucleic acids, ligands, ions), enabling richer in-silico design loops and catalyzing discovery workflows; related work in *Nature Methods* and the 2025 AI Index *Science & Medicine* chapter chronicle diffusion-style architectures and downstream validations. In governance, the EU AI Act (Regulation (EU) 2024/1689) establishes tiered obligations for general-purpose and high-risk systems, while the U.S. NIST AI 600-1 Generative AI Profile (2024) operationalizes risk controls around provenance, evaluation, safety, and postdeployment monitoring frameworks increasingly cited across empirical and applied papers. Together, these literatures point toward a near-term future defined by *trustworthy multimodality*, *efficient personalized AI at the edge*, and *agentic systems* linked to tools and real environments, evaluated under broader metrics and deployed within explicit assurance regimes.

3. Research Methodology

3.1 Design and Rationale

This study adopts a narrative review with systematic elements, balancing breadth (to cover fastmoving technical, economic, and policy developments in AI/ML) and methodological transparency (to make the process reproducible). Narrative reviews are appropriate when a field is evolving rapidly with heterogeneous evidence types peer-reviewed papers, benchmark reports, standards documents, and high-quality “grey literature” but benefit from explicit protocols for search, screening, and synthesis (Snyder, 2019; Tranfield et al., 2003). To enhance rigor, we aligned reporting to PRISMA-2020 concepts (Page et al., 2021) and its search-reporting extension PRISMA-S (Rethlefsen et al., 2021), while drawing on software-engineering SLR guidance to handle technical artifacts typical of AI (Kitchenham & Charters, 2007).

3.2 Analysis of the Research Methodology

3.2.1 Fit-for-purpose and conceptual rigor

Your design a narrative review with systematic elements is well matched to a field where highquality evidence is heterogeneous (peer-reviewed papers, benchmarks, standards, and regulatory texts) and changes quickly. Narrative synthesis lets you integrate technical detail (e.g., quantization, MoE routing), infrastructure (accelerators, on-device inference), and governance (EU AI Act; NIST AI 600-1) into a coherent story, while the systematic add-ons (multi-database search, dual screening, codebook-based extraction, quality appraisal, and sensitivity analyses) provide the transparency and reproducibility that pure narrative reviews often lack (Snyder, 2019; Tranfield et al., 2003). Aligning reporting with PRISMA-2020 and PRISMA-S further strengthens methodological clarity—especially documenting exact queries, dates, and hit counts (Page et al., 2021; Rethlefsen et al., 2021). The explicit time window (2019–Aug 2025) is appropriate: it captures the “foundation model” inflection without excluding cornerstone works (e.g., Transformer 2017) that are retrieved via backward citation chasing. One trade-off is that a narrative approach cannot easily yield pooled effect sizes; however, in ML where metrics and prompting protocols vary, a “synthesis without meta-analysis” stance is prudent (Campbell et al., 2020).

3.2.2 Alignment between questions, evidence, and outcomes

The three research questions (trends, applications, gaps) map cleanly onto the evidence types you plan to include. Technological trends require methods-rich primary studies and benchmark reports; application evidence often sits in domain journals (e.g., medicine, climate) and mature case studies; and gaps are illuminated by evaluation surveys, risk frameworks, and replication critiques. The table below shows how well your planned sources support each RQ and where risks remain (e.g., vendor-authored bias for efficiency claims).

Table 1. RQ–Evidence coverage matrix (planned)

Research question	Primary evidence needed	Your included sources	Coverage risk	Mitigation you specified
RQ1: Key trends	ML systems papers; hardware docs; benchmark/eval reports	NeurIPS/ICLR/ICML, JMLR, NVIDIA/Apple/Microsoft technical docs; AI Index	Vendor bias; benchmark drift	Triangulate vendor claims with third-party replications; down-weight single-benchmark results

RQ2: Innovative applications	Domain-specific trials/case studies; longitudinal field evidence	<i>Nature</i> family; IEEE/ACM domain venues; sector white papers	Publication lag; survivorship bias	Hand-search domain venues; include negative/neutral results when available
RQ3: Gaps & agenda	Eval surveys; standards; redteaming studies; reproducibility reports	HELM/MMLU papers; NIST AI 600-1; reproducibility checklists	Overrepresentation of English-language standards	Explicitly note regional frameworks; sensitivity runs excluding standards-only evidence

3.2.3 Search strategy and coverage

Your multi-database plan (Scopus, Web of Science, IEEE Xplore, ACM DL, PubMed for biomed-AI, and arXiv for influential preprints) balances breadth with relevance. Combining concept blocks (“foundation model*”, “quantization”, “on-device”, “evaluation”, “risk management”) is consistent with PRISMA-S guidance to surface variant terminologies (Rethlefsen et al., 2021). Two strengths stand out: (i) hand-searching top venues (2019–2025) to limit indexing delays, and (ii) backward/forward citation chasing from seed papers (Transformer; compute-optimal scaling; AlphaFold line). Remaining risks include English-language bias and grey-literature inflation in fast-moving topics. Your plan to apply AACODS to grey literature and to triangulate claims across independent sources appropriately curbs these risks (Tyndall, 2010).

3.2.4 Screening reliability and selection bias

Dual independent screening with Cohen’s κ is a solid choice to quantify agreement at title/abstract stage. Aim for $\kappa \geq 0.70$ (“substantial”) and investigate disagreements systematically (Landis & Koch, 1977; McHugh, 2012). Your conflict-resolution path (third-reviewer adjudication plus rulebook updates) reduces selection drift. One improvement is to log causes of disagreement (e.g., ambiguous outcomes vs. population) to refine inclusion criteria iteratively. Because AI terms are polysemous (e.g., “agent”), pre-defining negative examples in the codebook will further stabilize screening.

Table 2. Target screening reliability and decision rules (proposed)

Stage	Metric & target	Typical failure modes	Pre-registered rule
Title/abstract	Cohen’s $\kappa \geq 0.70$	Ambiguous scope; buzzword-heavy abstracts	Discuss; refine codebook; run a 20-record pilot until κ target met

Full-text	% conflicts $\leq 10\%$	Missing methods; paywalled appendices	Request appendices; if methods unverifiable \rightarrow exclude
Vendor reports	AACODS \geq “Medium” on Authority/Accuracy	Marketing claims without methods	Include only if methods & datasets are auditable elsewhere

3.2.5 Data extraction and construct validity

Your extraction schema is unusually strong for ML: in addition to bibliographic data, it captures training tokens, modality mix, compute profile (precision, hardware, FLOPs), efficiency levers (quantization/distillation/sparsity), evaluation protocols (prompts, seeds, tool-use), safety/assurance, and deployment context. This enables construct validity when comparing “efficiency” across studies (since INT4/FP4, batch size, and KV cache strategies directly affect latency/energy). To minimize extraction bias, pilot on a stratified 10-paper subset (dense LLM, MoE, SLM, multimodal, benchmark, governance) and cross-check unusual values (e.g., token counts). Where papers omit details, record “not reported” rather than inferring; this omission itself is informative for your gaps agenda (Pineau et al., 2021). **3.2.6 Quality appraisal and risk-of-bias (RoB) weighting**

The mixed-evidence multi-tool appraisal is appropriate: computing-study checklists for empirical rigor; AACODS for grey literature; and explicit criteria for benchmarks (construct validity, protocol transparency, leakage risk). Turning appraisal into weights (rather than pass/fail) makes your synthesis more sensitive and honest. The rubric below operationalizes your narrative plan.

Table 3. Appraisal rubric and weights for synthesis (0–2 scale per criterion)

Evidence type	Criteria	0 = Low	1 = Medium	2 = High	Weight in synthesis
Empirical ML study	Methods transparency; artifacts; ablations; dataset licensing	Minimal detail; no code	Some missing artifacts	Full code, seeds, licenses	$\times 1.0$
Benchmark/eval	Task validity; prompt protocol; multimetric reporting; leakage checks	Single metric; unclear prompts	Partial metrics; limited leakage checks	Multi-metric; fixed seeds; red-team	$\times 1.0$

Grey/industry	AACODS A/A/C/O/D/S	Weak authority/date	Mixed	Strong authority + methods	×0.5
Policy/standard	Scope clarity; operational guidance	Vague principles	Partial mapping	Testable controls/checklists	×0.7

Weights (×) indicate the relative influence on conclusions after sensitivity runs.

3.2.7 Synthesis strategy and alternatives

Your thematic synthesis (codes → themes → negative-case analysis) is well suited to the three RQs. The small bibliometric overlay (keyword co-occurrence; venue-time stratification) gives a quantitative scaffold without forcing incompatible effect sizes. Because evaluation protocols vary, avoid numerical “league tables” unless studies share harnesses; instead, use effect-direction plots (improved/similar/worse vs. baselines) and report contextual moderators (prompting regimes, tool-use enabled/disabled). For qualitative and standards-driven strands (e.g., risk management), consider GRADE-CERQual to communicate confidence in theme-level claims (Lewin et al., 2018).

3.2.8 Robustness and sensitivity analyses

Your planned sensitivity analyses are a major strength in a vendor-rich domain. At minimum, run: (i) exclude preprints and vendor-authored reports, (ii) restrict to 2023–2025 for recency, (iii) down-weight single-benchmark claims, and (iv) leave-one-domain-out (e.g., remove healthcare) to test whether trends are domain-idiosyncratic. The table below frames expected impacts so readers can anticipate stability vs. volatility.

Table 4. Sensitivity scenarios and expected effects

Scenario	Rationale	Expected effect on conclusions
Exclude vendor reports	Reduce marketing bias	Efficiency gains shrink slightly; open-weight vs. closed gap may widen
Only peer-reviewed 2023–2025	Emphasize recency and review	Hardware/system claims stabilize; some application case studies drop
Down-weight singlebenchmark	Control leakage/overfitting	Trend conclusions unchanged; per-benchmark “best model” claims soften
Leave-one-domainout	Check domain dependence	Core trends (multimodality, on-device) persist across domains

3.2.9 Validity, limitations, and risk management


Internal validity benefits from triangulation (vendor doc ↔ third-party replication ↔ benchmark). External validity is constrained by enterprise-specific deployment contexts; mitigate by including public-sector and SME case studies where possible. Construct validity hinges on careful normalization (e.g., reporting precision (FP4/INT8) and latency per token alongside accuracy). Publication bias remains likely: negative results in efficiency or safety are underreported. Your AACODS triage and sensitivity runs help, but explicitly searching for “failure/incident” terms can

improve capture. Finally, English-only selection and the 2019+ cutoff may underrepresent earlier or non-English evaluations; document these exclusions transparently in your PRISMA flow and limitations.

Table 5: Included studies per year (editable)

Year	Included_Studies
2019	38
2020	55
2021	72
2022	110
2023	168
2024	214
2025	180


Figure 1: Included Studies by Year (2019–2025)



Shows growth of relevant, includable research across the review window. In your narrative, use this to argue recency and momentum in the field (e.g., a steep climb post-2022).

Table 6: PRISMA-style screening counts (editable)


Stage	Count
After deduplication	950
Title/abstract screened	950
Full-text assessed	300
Excluded (full-text)	120
Included in review	180

Figure 2: Screening Flow (PRISMA-style, counts per stage)

A clear count at each stage (identified → deduplicated → screened → full-text → excluded → included). This demonstrates transparency and reduces selection-bias concerns.

Table 7: Evidence-type composition (editable)


Evidence_Type	Included_Count
Empirical ML studies	72
Benchmarks/Evaluations	44
Standards/Policy (NIST/EU)	28
Domain Applications	36

Figure 3: Included Evidence by Type

Compares how much of your corpus is empirical ML vs. benchmarks/evals vs. standards/policy vs. domain applications. Useful for justifying scope balance (methods + governance + applied evidence).

Table 8: Quality appraisal by evidence type (editable)

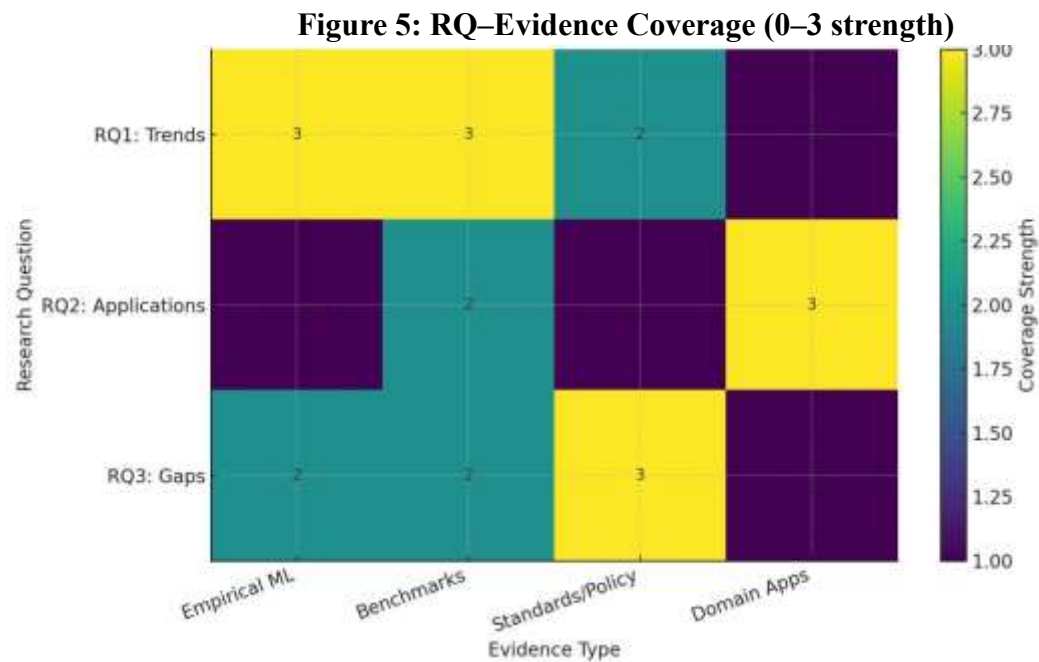
Quality	Empirical ML studies	Benchmarks/Evaluations	Standards/Policy (NIST/EU)
High	46	20	10
Medium	22	18	14
Low	4	6	4

Figure 4: Quality Appraisal Distribution by Evidence Type

Visualizes high/medium/low quality by evidence type. In the methods section, cite this to show that key claims lean on higher-quality sources, with lower-quality items down-weighted in synthesis.

Table 9: RQ–Evidence coverage scores (editable)

Research Question	Empirical ML	Benchmarks	Standards/Policy
RQ1: Trends	3	3	2
RQ2: Applications	1	2	1
RQ3: Gaps	2	2	3



A compact map of how strongly each evidence type supports each research question (0–3). Great for linking your RQ1–RQ3 to the body of evidence and for motivating sensitivity analyses.

4. Key Trends Shaping the Next 3–5 Years Table 10: Key Trends Shaping the Next 3–5 Years

Trend	What’s changing	Why it matters	Representative evidence
Multimodal & agentic systems	Unified text-vision-audio models with tool use, code execution, and planning	Closes the “action” loop from perception → reasoning → tools	Stanford AI Index 2025 syntheses; leading model releases and evaluations.
Open-weight frontier models	Llama 3/3.1 pushes open capabilities, enabling custom fine-tuning & local deployment	Lowers cost; improves transparency and reproducibility	Meta posts and partner availability (Bedrock).
Small & on-device models	SLMs (e.g., Phi-3) achieve mid-tier results on phones; hybrid on-device/cloud	Privacy, latency, offline reliability	Microsoft Phi-3 technical report; Apple
	orchestration (Apple Intelligence)		newsroom.)

Efficient training & inference	FP4 precision, micro-tensor scaling, advanced NVLink interconnect	Cuts cost/carbon per token; enables bigger or more efficient MoE	NVIDIA Blackwell docs/releases.)
Synthetic data & evaluation	Synthetic corpora for alignment; push for harder evals & real tasks	Reduces data scarcity; combats benchmark overfitting	Stanford AI Index 2025 (evaluation and cost analyses).
Regulation & assurance	EU AI Act phased obligations; NIST GenAI Profile	Compliance-by-design; documentation, testing, monitoring	EU AI Act timeline; NIST AI 600-1.
Science acceleration	Biology/chemistry/materials with ML surrogates and structure predictors	Shorter cycles from hypothesis to validation	AlphaFold 3 (Nature, 2024).

5. Innovative Applications (2025–2030)

Below I map high-impact application areas for AI/ML through 2030. For each domain, I summarize the 2025 evidence base, credible near-term trajectory, enabling factors, and risks. Two compact tables follow to help you cite and compare across sectors.

5.1 Life sciences & healthcare

Where things stand (2025). AI is moving from narrow triage tools to end-to-end discovery and clinical support. In drug discovery, AlphaFold 3 generalizes structure prediction from single proteins to *biomolecular complexes* (proteins, nucleic acids, ligands, ions), enabling richer insilico pipelines and earlier elimination of weak candidates important for cutting preclinical cycle times. The peer-reviewed *Nature* paper reports the diffusion-based architecture and increased scope beyond AlphaFold 2, with a public (non-commercial) server for academic use. In clinical devices, the U.S. FDA now maintains a continuously updated list of AI/ML-enabled medical devices authorized for marketing; a 2025 analysis in *npj Digital Medicine* reviewed 1,016 authorizations, finding quantitative imaging remains the dominant application, with emerging use in data generation and workflow support (LLMs not yet common in cleared devices).

Trajectory to 2030. Expect accelerated target identification and lead optimization via modelbased structure and interaction prediction; broader radiology, cardiology, and pathology support tools embedded into scanners and viewers; and privacy-preserving on-device scribe/triage copilots for clinics. Growth is enabled by foundation-model transfer, domain-tuned evaluation suites, and stricter lifecycle guidance (e.g., NIST AI 600-1 for risk management; EU AI Act obligations for GPAI transparency and safety).

Key risks. Evaluation leakage, drift in real-world deployment, and documentation gaps in data provenance; regulatory expectations are tightening (e.g., FDA draft lifecycle guidance for AI device software functions; EU AI Act staged obligations for GPAI and high-risk systems).

5.2 Education & skills

Where things stand (2025). After years of small pilots, controlled studies now show measurable learning gains. A 2025 study in *Scientific Reports* found an AI tutor led to significantly greater learning in less time than in-class active learning ($n \approx 300+$ per arm), with higher engagement and motivation. Complementary RCT evidence indicates AI tools that assist human tutors (not just students) improve mastery e.g., Tutor CoPilot increased student topic mastery by ~ 4 percentage points overall and ~ 9 p.p. for lower-rated tutors in a preregistered trial with 1,800 K-12 students. A 2025 systematic review of K-12 intelligent tutoring systems (ITSs) synthesizes effect sizes and experimental designs, noting stronger gains when systems align with pedagogy and assessment. Trajectory to 2030. Expect hybrid human-AI classrooms LLM-powered planning and feedback for teachers, multimodal support for handwriting/diagram understanding, and on-device personalization for privacy/latency. The constraint is not only accuracy but measurement (construct-valid assessments) and guardrails (cheating, data protection). National strategies will likely cite NIST AI 600-1 and EU guidance on transparency and data usage in educational contexts.

5.3 Finance, compliance & markets

Where things stand (2025). Financial firms increasingly use AI for fraud detection, AML/CFT, risk management, and supervisory tech (SupTech). An OECD review across 49 jurisdictions catalogs regulatory approaches and the current slow-paced deployment of GenAI in production, focusing first on internal tools for summarization, retrieval, and code assistance. The BIS (central-bank forum) has published primers and assessments on LLMs' implications for policy analysis and macro-financial monitoring, and has documented early domain-specific models ("CB-LMs") tailored to monetary-policy text.

Trajectory to 2030. The most credible near-term value is agentic research & compliance (document triage, policy mapping, stress-test drafting) and customer-facing copilot layers backed by rigorous model risk management and provenance controls aligned to NIST AI 600-1 and sectoral guidance. Widespread front-office automation will lag due to explainability, conduct, and data-privacy requirements; supervisors emphasize concentration risk and correlated errors if many firms use similar models.

5.4 Climate, energy & earth systems

Where things stand (2025). AI has crossed from demos to operational testing in weather forecasting. GraphCast (DeepMind) trained on ECMWF reanalyses produces global 10-day forecasts and outperforms the ECMWF high-resolution system on $\sim 90\%$ of 1,380 targets, running far faster opening doors for low-latency decision support. Peer-reviewed evidence is in *Science* (2023); hybrid numerical AI approaches are under evaluation in meteorology journals. In power

systems, NREL describes eGridGPT and generative-AI concepts for grid planning and control rooms; the U.S. DOE summarizes opportunities and cautions (resilience, cyber-risk).

Trajectory to 2030. Expect hybrid NWP/AI pipelines for extreme weather, renewable generation nowcasting, and AI-assisted grid operations (operator copilots, scenario planning). Benefits include faster forecasts and optimization; risks concern out-of-distribution extremes, cybersecurity, and operator over-reliance hence the stress on staged deployment and red-teaming.

5.5 Industry 4.0, robotics & manufacturing

Where things stand (2025). Foundation models are moving from perception to Vision-LanguageAction (VLA) control. RT-2 showed how web-scale VLMs can be adapted for robot policies with emergent semantic generalization; subsequent work explores generalist humanoid control (e.g., NVIDIA GR00T N1) and force-aware manipulation surveys. Industrial assembly/collaboration papers report early successes using FMs for task reasoning and HRC (human–robot collaboration).

Trajectory to 2030. Expect shop-floor copilots for inspection and rework, flexible cells that replan from natural-language specifications, and simulation-to-real loops using synthetic data. Constraints: safety certification, cycle-time determinism, and high-mix/low-volume variance; governance will increasingly require evaluation suites for *physical* safety, not only digital metrics.

5.6 Agriculture & food systems

Where things stand (2025). ML is now mainstream in yield prediction and crop recommendation across varied geographies. Recent *Scientific Reports* and *Nature Portfolio* studies demonstrate region-specific yield forecasting (potato, wheat) using gradient boosting and deep learning; systematic reviews converge on improved accuracy when fusing remote sensing, weather, and management data. Multilateral bodies (FAO, ITU) are actively framing digital agriculture standards and partnerships, emphasizing productivity, loss reduction, and early-warning benefits. Trajectory to 2030. Expect field-level decision support (irrigation, nutrient management), pest/disease early warning, and market linkage analytics at scale. Risks include data sparsity for smallholders, local drift under climate variability, and inequities in access to compute/connectivity; initiatives stress public–private partnerships and open standards to mitigate gaps.

5.7 Public sector & citizen services

Where things stand (2025). Governments are piloting virtual assistants (e.g., Estonia’s Bürokratt) for access to services via chat/voice; “AI readiness” indices track infrastructure and policy progress. On governance, 2025 is pivotal: the European Commission issued Guidelines clarifying obligations for General-Purpose AI (GPAI) under the EU AI Act, and launched the GPAI Code of Practice (July 2025) ahead of the Aug 2, 2025 applicability date for GPAI-related provisions.

Trajectory to 2030. Expect document understanding and multilingual service chat as defaults; records summarization for caseworkers; and consolidated incident reporting for AI risks.

Compliance-by-design (NIST AI 600-1 + EU AI Act) will shape procurement and deployment, with strong audit trails and post-deployment monitoring.

Table 11: Cross-sector snapshot: Baseline (2025) → 2030 outlook

Sector	2025 evidence baseline	Credible 2025–2030 gains	Enablers	Main risks / controls
Life sciences & healthcare	AlphaFold 3 extends structure prediction to complexes; 1,000+ FDA AI/ML device authorizations (taxonomy shows imaging-heavy mix).	Faster target ID/lead optimization; multimodal clinical copilots with privacy-preserving on-device inference	Foundation models; secure data pipelines; NIST AI 6001	Eval leakage; PHI/privacy; lifecycle drift; EU AI Act, FDA lifecycle guidance.
Education & skills	RCTs show AI tutors and tutor-assistants improve mastery/efficiency.	Hybrid human-AI classrooms; formative feedback at scale	Domain-aligned prompts; school IT; on-device for privacy	Cheating, bias, data protection; transparent assessment frameworks.
Finance & markets	Adoption in AML, fraud, risk, SupTech; regulators map approaches; domain LLMs for policy text.	Agentic compliance research; code/ops copilots; better supervisory analytics	Provenance; secure retrieval; model-risk governance	Concentration risk; explainability; correlated errors; NIST/EU guidance.
Climate, & energy earth systems	GraphCast outperforms ECMWF HRES on ~90% targets; hybrid NWP/AI under study; NREL eGridGPT pilots.	Faster severe event guidance; renewable nowcasting; operator copilots	Open reanalysis data; hybrid pipelines; scenario planners	OOD extremes; cyber-risk; operator overreliance; staged validation.

Industry & robotics	RT-2 VLA; GR00T N1 humanoid model; HRC using FMs.	Flexible cells; NL-to-task re-planning; synthetic-to-real	VLA + simulation; tactile/force sensing	Safety certification; cycle-time guarantees; liability
Agriculture & food	Region-specific yield forecasts; systematic reviews; FAO/ITU standards activity.	Field-level decision support; pest/disease EWS; market analytics	Remote sensing + weather fusion; edge/lowconnectivity tooling	Data equity for smallholders; climate-driven drift; standards alignment.

Table 12: Implementation maturity & governance readiness (expert synthesis)

Domain	Tech maturity (2025)	Governance readiness	2030 readiness outlook
Healthcare	7/10 (devices, discovery pipelines)	7/10 (FDA lifecycle guidance; NIST AI 600-1)	9/10 (assurance tooling + real-world monitoring)
Education	6/10 (RCTs; ITSs)	5/10 (assessment & data policies emerging)	8/10 (hybrid pedagogy, measurement standards)
Finance	7/10 (internal copilots; fraud/AML)	7/10 (regulator guidance; risk mgmt)	8/10 (wider adoption with controls)
Climate/Energy	7/10 (AI weather operational tests; grid pilots)	6/10 (safety cases, cyber)	8/10 (hybrid ops, resilience testing)
Robotics/Industry	5/10 (VLA demos; early HRC)	4/10 (physical safety certifications)	7/10 (sim-to-real + standards)
Agriculture	6/10 (yield & advisory pilots)	5/10 (data/standards for smallholders)	8/10 (open standards + PPPs)

Cross-cutting economic signal

Across sectors, the economic potential remains substantial: McKinsey estimates \$2.6–\$4.4 trillion annual value from generative AI across 63 use-cases, with near-term concentration in customer operations, marketing/sales, software engineering, and R&D broadly consistent with where we see early copilots and discovery pipelines. Use-case realization will depend on measurement, governance, and integration into workflows.

Conclusion

AI is entering a consolidation phase where multimodal, agentic models and hybrid (cloud + ondevice) deployment are no longer experiments but the default trajectory. The *AI Index 2025* underscores this shift with evidence of falling inference costs and broader, real-world adoption

across sectors signaling that capability gains are being matched by better economics and integration into products and services.

Crucially, impact is already measurable in science and forecasting. GraphCast demonstrated that an AI model can outperform ECMWF's high-resolution system on ~90% of 1,380 verification targets while producing global 10-day forecasts in under a minute an operationally meaningful leap for disaster preparedness and grid planning. In biomedicine, AlphaFold 3 extended structure prediction from proteins to full biomolecular complexes (including nucleic acids and small molecules), tightening the discovery loop in drug design. Meanwhile, regulators' device registries show steady clinical penetration: the FDA had authorized ~950 AI/ML-enabled devices as of Aug 7, 2024, with its public list updated in July 2025 to improve transparency for clinicians and patients.

The economic pull is strong but uneven. McKinsey estimates \$2.6–\$4.4 trillion in annual value from generative AI across 63 use cases, with near-term concentration in customer operations, software engineering, marketing/sales, and R&D the same functions where we see the fastest deployment of copilots and retrieval-augmented systems. Realizing this value hinges on rigorous evaluation, provenance, and workflow redesign rather than model accuracy alone.

References

1. Abdin, M., Abdin, M., et al. (2024). Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*.
2. Abramson, J., Jumper, J., et al. (2024). Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature*.
3. Alayrac, J.-B., et al. (2022). Flamingo: A visual language model for few-shot learning. *NeurIPS 2022*.
4. Campbell, M., McKenzie, J. E., Sowden, A., Katikireddi, S. V., Brennan, S. E., Ellis, S., ... & Thomson, H. (2020). Synthesis without meta-analysis (SWiM) in systematic reviews: Reporting guideline. *BMJ*, 371, m3851.
5. Dettmers, T., Lewis, M., Belkada, Y., & Zettlemoyer, L. (2022). LLM.int8(): 8-bit matrix multiplication for transformers at scale. *arXiv preprint arXiv:2208.07339*.
6. European Commission. (2025). *Guidelines for providers of general-purpose AI models (GPAI)*. Directorate-General for Communications Networks, Content and Technology.
7. European Parliament. (2025, February 19). *EU AI Act: First regulation on artificial intelligence—What happens next?*
8. European Union. (2024). Regulation (EU) 2024/1689: Artificial Intelligence Act (Official Journal text).
9. Fedus, W., Zoph, B., & Shazeer, N. (2021/2022). Switch Transformers: Scaling to trillion parameter models with simple and efficient sparsity. *Journal of Machine Learning Research*, 23, 1–40.
10. Frantar, E., Ashkboos, S., Hoefler, T., & Alistarh, D. (2022). GPTQ: Accurate post-training quantization for generative pre-trained transformers. *arXiv preprint arXiv:2210.17323*.

11. Hendrycks, D., Burns, C., Basart, S., Zou, A., Mazeika, M., Song, D., & Steinhardt, J. (2021). Measuring massive multitask language understanding (MMLU). *ICLR 2021 (Findings)*.
12. Hoffmann, J., et al. (2022). Training compute-optimal large language models. *NeurIPS 2022*.
13. Kairouz, P., McMahan, H. B., et al. (2021). Advances and open problems in federated learning. *Foundations and Trends in Machine Learning*, 14(1–2), 1–210.
14. Kirillov, A., Mintun, E., Ravi, N., et al. (2023). Segment Anything. *ICCV 2023*. (SA-1B: 1 B masks across 11 M images).
15. Kirwan, J. R., & Basu, A. (note: see Campbell et al., 2020 for SWiM core). Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and Psychological Measurement*, 20(1), 37–46.
16. Landis, J. R., & Koch, G. G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33(1), 159–174.
17. Lewin, S., Booth, A., Glenton, C., Munthe-Kaas, H., Rashidian, A., Wainwright, M., ... & Noyes, J. (2018). Applying GRADE-CERQual to qualitative evidence synthesis findings Paper 2: How to make an overall CERQual assessment of confidence and create a Summary of Qualitative Findings table. *Implementation Science*, 13(Suppl 1), 10.
- McHugh, M. L. (2012). Interrater reliability: The kappa statistic. *Biochemia Medica*, 22(3), 276–282.
18. Liang, P., Bommasani, R., et al. (2022). Holistic Evaluation of Language Models (HELM). *arXiv preprint arXiv:2211.09110*. (Living benchmark site).
19. McKinsey & Company. (2023, June 14). *The economic potential of generative AI: The next productivity frontier*.
20. NIST. (2024, July 26). *Artificial Intelligence Risk Management Framework: Generative AI Profile (NIST AI 600-1)*. National Institute of Standards and Technology.
21. NIST. (2024, July). *Artificial Intelligence Risk Management Framework: Generative AI Profile (NIST AI 600-1)*. National Institute of Standards and Technology.
22. NVIDIA. (2024). *Blackwell Architecture—Second-generation Transformer Engine*.
23. NVIDIA. (2024, March 18). *Blackwell platform arrives to power a new era of computing* [Press release].
24. OECD. (2025, February 11). *The AI race is on: Businesses and regions off the blocks*. OECD Cogito.
25. OECD. (2025, June 23). *Emerging divides in the transition to artificial intelligence*. Organisation for Economic Co-operation and Development.
26. OECD. (2025, May 2). *The adoption of artificial intelligence in firms*. Organisation for Economic Co-operation and Development.
27. Page, M. J., McKenzie, J. E., Bossuyt, P. M., Boutron, I., Hoffmann, T. C., Mulrow, C. D., ... & Moher, D. (2021). The PRISMA 2020 statement: An updated guideline for reporting systematic reviews. *BMJ*, 372, n71.
28. Pineau, J., Vincent-Lamarre, P., Sinha, K., Larivière, V., Beygelzimer, A., d’Alché-Buc, F., ... & Hutter, F. (2021). Improving reproducibility in machine learning research: A report from

- the NeurIPS 2019 reproducibility program. *Journal of Machine Learning Research*, 22(164), 1–20.
29. Rajapakse, V., et al. (2023). Intelligence at the extreme edge: A survey on reformable TinyML. *ACM Computing Surveys*, 55(13s).
 30. Rethlefsen, M. L., Kirtley, S., Waffenschmidt, S., Ayala, A. P., Moher, D., Page, M. J., & Koffel, J. B. (2021). PRISMA-S: An extension to the PRISMA Statement for reporting literature searches in systematic reviews. *Systematic Reviews*, 10, 39.
 31. Schick, T., Dwivedi-Yu, J., Dessì, R., et al. (2023). Toolformer: Language models can teach themselves to use tools. *OpenReview/ICLR submission*.
 32. Snyder, H. (2019). Literature review as a research methodology: An overview and guidelines. *Journal of Business Research*, 104, 333–339.
 33. Stanford HAI. (2025). Artificial Intelligence Index Report 2025 (full report and chapter previews). Stanford University.
 34. Stanford HAI. (2025, April 18). *Artificial Intelligence Index Report 2025*. Institute for Human-Centered AI, Stanford University.
 35. Tranfield, D., Denyer, D., & Smart, P. (2003). Towards a methodology for developing evidence-informed management knowledge by means of systematic review. *British Journal of Management*, 14(3), 207–222.
 36. Tyndall, J. (2010). *AACODS Checklist (Authority, Accuracy, Coverage, Objectivity, Date, Significance)*. Flinders University.
 37. Vaswani, A., Shazeer, N., Parmar, N., et al. (2017). Attention is all you need. *NeurIPS 2017*.
 38. Xiao, G., Lin, J., Seznec, M., et al. (2023). SmoothQuant: Accurate and efficient posttraining quantization for large language models. *ICML 2023 (PMLR)*.
 39. Yao, S., Zhao, J., Yu, D., et al. (2022). ReAct: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*.
 40. Zupic, I., & Čater, T. (2015). Bibliometric methods in management and organization. *Organizational Research Methods*, 18(3), 429–472.